

## Laboratory 2

In this laboratory we will introduce the IRIS dataset. We will see how we can load and visualize the dataset, and compute some statistics of the data.

### The IRIS Dataset

The dataset was introduced in Fisher, R.A. “The use of multiple measurements in taxonomic problems”, Annual Eugenics, 7, Part II, 179-188 (1936). The dataset contains information on 150 samples (instances) of iris flowers belonging to 3 different families (classes): iris setosa, iris versicolor and iris virginica. There are 50 samples for each class. For each sample, the dataset provides 4 attributes: sepal length (cm), sepal width (cm), petal length (cm), petal width (cm).

### Loading the dataset

We want to load the dataset into a numpy 2-dimensional array of shape  $(4 \times 150)$ . Each row of the data matrix will correspond to an attribute, whereas each column will represent a sample. We will also construct a 1-dimensional array of class labels. Iris setosa will be indicated with value 0, iris versicolor with value 1 and iris virginica with value 2.

The dataset is provided as a comma-separated values (csv) file, where each line represents a data point. Each line contains the four attributes and the family name, for example:

5.1,3.5,1.4,0.2,Iris-setosa

The dataset can be found in `Solution/iris.csv`. Write a `load` function that, given the dataset filename, returns the  $(4 \times 150)$  data matrix, and the corresponding 1-dimensional array of size 150 containing the class label. Suggestion: for each line, create a  $4 \times 1$  vector with the attribute values. Store all vectors in a list, and then concatenate the vectors.

### Visualizing the dataset

We want to visualize the distribution of the different attributes for the different classes.


- For each attribute, plot the corresponding histogram for each class. You can use `matplotlib.pyplot.hist`, providing a 1-dimensional numpy array with the target data. Try with different number of bins.

You can normalize the histogram setting the parameter `density = True`. To create a new figure, you can use `matplotlib.pyplot.figure`. To visualize the current figure(s), use `matplotlib.pyplot.show`.

NOTE: `matplotlib.pyplot.show` shows the current image(s) in different windows, and blocks the execution until the windows are closed.

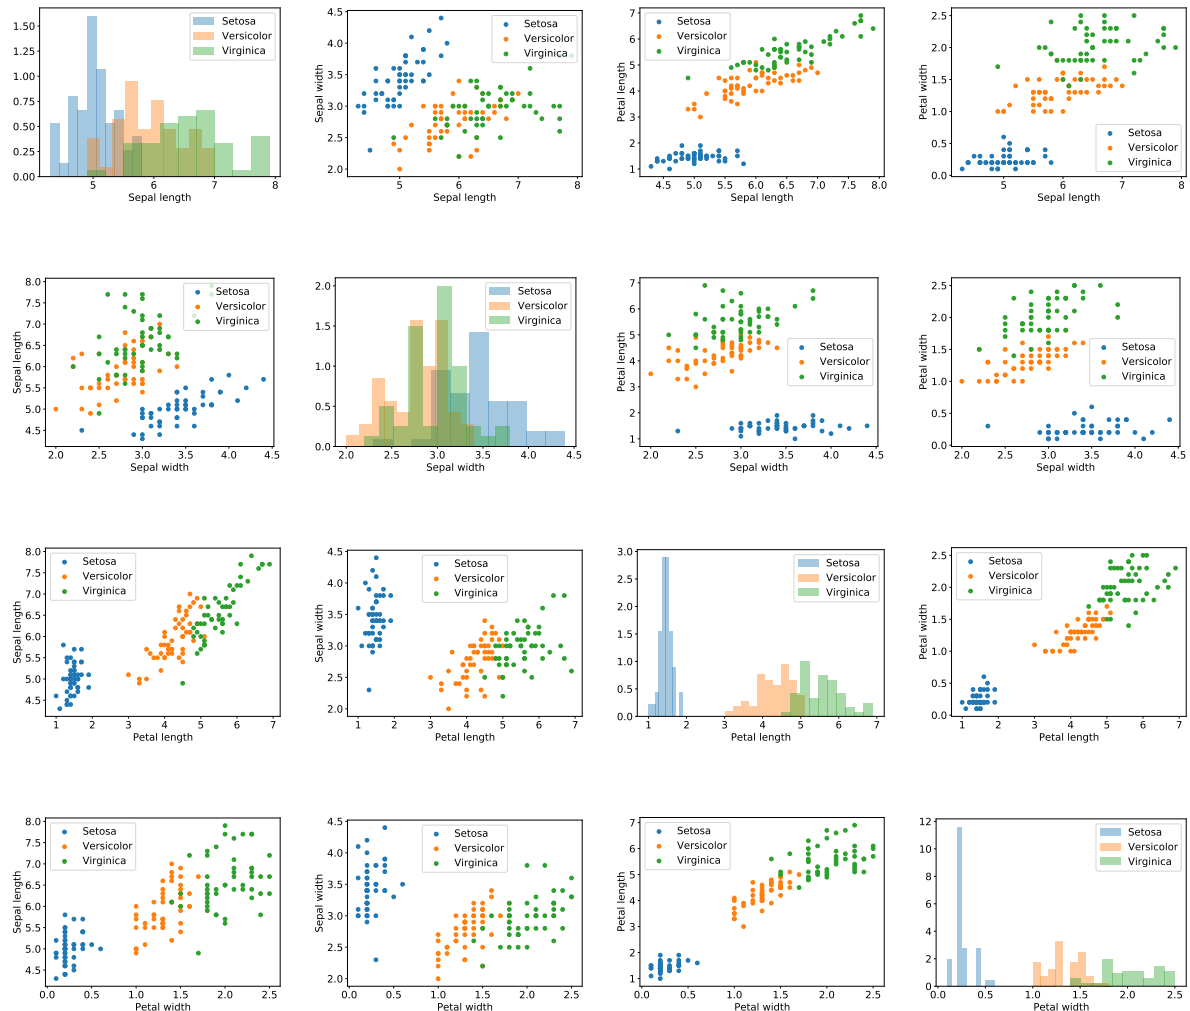
Suggestion: start extracting from the data matrix the parts corresponding to the different classes, then proceed with the plots. If `D` is the data matrix, and `L` is the label vector (1-dimensional), we can build a mask for each class and use it to filter the columns of `D`:

```
M0 = (L == 0)
D0 = D[:, M0] # or equivalently D0 = D[:, L==0]
```



We can observe that there is large overlap for the first two attributes, whereas values for the third and fourth attributes of iris-setosa are well separated from those of the other two flower families

- We now consider pairs of values. Visualize the scatter plots of the different attribute pairs for each class.



## Statistics

We now look at how to compute simple statistics and transformations of the data.

### Data mean

We can compute the dataset mean using a **for** loop:

```
mu = 0
for i in range(D.shape[1]):
    mu = mu + D[:, i:i+1]

mu = mu / float(D.shape[1])
```

The **for** loop is, in general, slow. **Numpy** allows computing the mean of an array through the method **.mean**. The method allows specifying an axis — for 2-D arrays, **axis = 0** allows computing the mean over rows, whereas **axis = 1** allows computing the means over columns. We can thus compute the dataset mean as

```
mu = D.mean(1)
```

Pay attention to the shape of `mu`: it's a 1-D array.

We now exploit broadcasting to *center the data*, i.e. to remove the mean from all points:

```
DC = D - D.mean(1).reshape((D.shape[0], 1))
```

Notice that we want the mean to be a column vector, so we first reshape it.

*Suggestion:* we will often have to reshape 1-D vectors as column or row vectors. Write a function `mcol` and a function `mrow` that implements the reshaping

You can try plotting again the centered data.